



Project #1: Artificial Neural Networks

Amir Sadovnik
COSC 525: Deep Learning (Spring 2021)

1 Overview

In this project you will be implementing a simple artificial neural network (ANN) library. I will ask you to implement it in a specific object oriented manner. Although this will make the library less efficient it will allow you to gain a better understanding of how ANN's predict and learn.

Make sure to use your own solution and your own code.

2 Problem Description

You are tasked with writing the following functions:

1. Write a `Neuron` class with the following properties:
 - (a) Should be initialized with an activation function, the number of inputs, the learning rate, and possibly a vector of weights (if not set to random).
 - (b) Needs to store the the input, output and partial derivative for each of its weights for back-propagation.
 - (c) Should have an `activate` method which given a value returns its value after activation (depending on the activation function used).
 - (d) Should have a `calculate` method which given an input calculates the output.
 - (e) Should have a `activationderivative` method which returns the derivative of the activation function with respect to the net.
 - (f) Should have a `calcpartialderivative` method which calculates the partial derivative with respect to each weight (using the `derivative` method) and returns the vector $w \times \delta$.
 - (g) Should have a `updateweights` which actually changes the weights using the partial derivative values (and the learning rate).

2. Write a `FullyConnectedLayer` class.

- (a) Should be initialized with the number of neurons in the layer, the activation function for all the neurons in the layer, the number of inputs, the learning rate, and possibly a vector of weights (if not set to random).
- (b) Should have a `calculate` method which given an input calculates the output of all the neurons in the layer.
- (c) Should have a `calculatewdeltas` which given the $\sum w \times \delta$ from the next layer, goes through all the neurons in the layer to calculate their partial derivative (using the `calcpartialderivative` method), updates all the weights (using the `updateweights` method) and returns its own $\sum w \times \delta$.

3. Write a `NeuralNetwork` class.

- (a) Should be initialized with the number of layers, number of neurons in each layer, a vector of activation functions for each layer, the number of inputs to the network, the loss function used, the learning rate, and possibly a vector of weights (if not set to random).
- (b) Should have a `calculate` method which given an input calculates the output of the network.
- (c) Should have a `calculateloss` method which given an input and desired output calculates the loss.
- (d) Should have a `lossderivative` method which returns the value of the loss derivative (depending on the loss).
- (e) Should have a `train` method, which given a single input and desired output takes one step of gradient descent. This method should first do a forward pass (give its own method
- (f) Any additional methods/instance variables needed for back-propagation `calculate`, then calculate the derivative of the loss (using `lossderivative`) and finally go through the layers backwards calling `calculatewdeltas` for each and passing $\sum w \times \delta$. as it goes. Note that the first time `calculatewdeltas` is called it should be passed the derivative loss.

4. Your `main` method should take the following command line variables to show the following:

- (a) If given `example`, simply run a single step of back-propagation using the example we did in class. Print out the weights after the 1-step update to show that you have gotten the same results as we did in class.
- (b) If given `and` should train a single perceptron to produce the output of the “and” logic gate. Train it for enough steps so that it converges. Show the results of the prediction for all 4 possible inputs.

- (c) If given `xor` should train two types of networks to produce the output of the “xor” logic gate. One is a single perceptron, while the other adds a hidden layer. Train it for enough steps so that it converges. Show the results of the prediction for all 4 possible inputs for each of them.

3 Additional Information

You must do so under the following constraints:

1. Your library should support at least two activation functions: logistic and linear
2. Your library should support at least two loss functions: square error and binary cross entropy loss
3. You must use Python3 with only the numpy and sys libraries allowed
4. Make sure to comment your code
5. **Work plan tip:** Start with implementing a feed-forward network. First implement the neuron methods for feed forward, then the `fullyconnectedlayer` for feed forward, and finally `neuralnetwork` for feed forward. For each of them you can do some manual calculations to ensure you are getting the correct results. Only then start working on back-propagation. Do so in the same order checking manually along the way.

4 Report

You should submit a short PDF report with the following (if you do not have anything to add for a section, simply put the section title and then state there is nothing to add):

1. A short introduction to the problem.
2. Assumptions/choices you have made.
3. Problems/Issues you were not able to solve.
4. How to run your code (if there is any difference from what is stated)
5. Show a graph of the loss dropping as a function of learning for different learning rates. Try to find learning rates that are too high/low. Describe what you see in the graph in words.

5 Submission

You are required to submit one zip file with the code and your report.